

Agents do not just answer. They choose, call tools, and act.

Holmes audits high-risk agent decisions before deployment and after incidents. It converts an agent action into a signed evidence artifact: what information pushed the action, how close it was to a different action, and whether policy gates should block release.

TOOL RISK

Tiered

payment, deletion, email, external API

GATE

Pass/Block

policy result with reasons

ARTIFACT

Signed

hashes, replay metadata, thresholds

What Holmes tests

Action margin, close-to-flip prompts, contradictory instruction failures, tool-call counterfactuals, and model-update drift across approved benchmark prompts.

What buyers get

A release gate for agentic workflows: safe actions pass, risky actions create review tickets, and incidents can be reconstructed from preserved evidence.

Claim boundary

For black-box agents Holmes gives behavioral evidence. For open-weight agents it can escalate to mechanistic circuit tests when weights are available.