

A bounded internal circuit claim, not a generic explanation.

For open-weight models, Holmes can move beyond surface attribution. It reads the exact decision margin, finds internal layer-token states, intervenes on activations, and certifies only when the selected circuit passes sufficiency, necessity, minimality, and negative controls.

READOUT

Logit

target vs alternative margin

CIRCUIT

Patched

internal states are intervened on

CONTROLS

Required

random and wrong-target tests fail

Certification gates

Clean/corrupt contrast, decision readout, layer timing, activation patching, sufficiency, necessity, minimality, and negative controls.

Allowed claim

For this model checkpoint, prompt family, target decision, and decision metric, Holmes found a circuit that was necessary and sufficient within the tested domain.

Not certified

Global model behavior, every future prompt, human-like reasoning, truth of the answer, or legal causation without independent validation.